# Introduction to Kyoto Products

# **Kyoto Products**

- Kyoto Cabinet
  - a lightweight database library
    - a straightforward implementation of DBM
- Kyoto Tycoon
  - a lightweight database server
    - a persistent cache server based on Kyoto Cabinet



- lightweight database library -

### Features

- straightforward implementation of DBM
  - process-embedded database
  - persistent associative array = key-value storage
  - successor of QDBM, and sibling of Tokyo Cabinet
    - e.g.) DBM, NDBM, GDBM, TDB, CDB, Berkeley DB
  - C++03 (with TR1) and C++0x porable
    - supports Linux, FreeBSD, Mac OS X, Solaris, Windows
- high performance
  - insert 1M records / 0.8 sec = 1,250,000 QPS
  - search 1M records / 0.7 sec = 1,428,571 QPS

### high concurrency

- multi-thread safe
- read/write locking by records, user-land locking by CAS

### high scalability

- hash and B+tree structure = O(1) and O(log N)
- no actual limit size of a database file (up to 8 exa bytes)

### transaction

- write ahead logging, shadow paging
- ACID properties

### various database types

- storage selection: on-memory, single file, multiple files in a directory
- algorithm selection: hash table, B+ tree

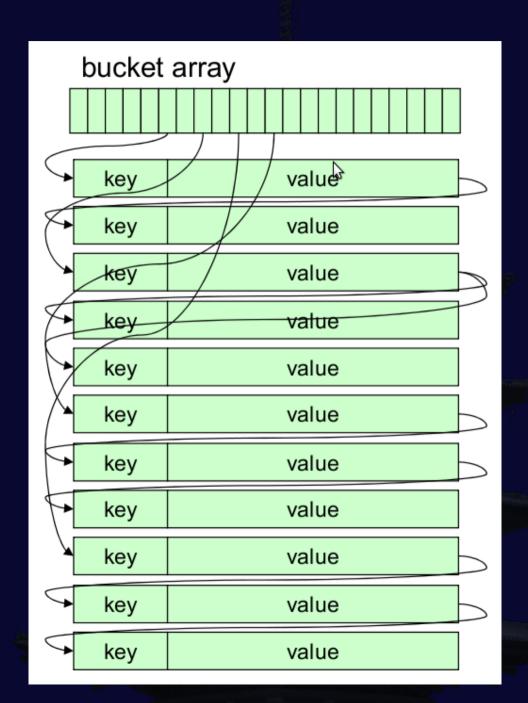
### other language bindings

· C, Java, Python, Ruby, Perl, Lua, and so on

# HashDB: File Hash Database

### static hashing

- O(1) time complexity
- jilted dynamic hashing for simplicity and performance
- separate chaining
  - binary search tree
  - balances by the second hash
- free block pool
  - best fit allocation
  - dynamic defragmentation
- combines mmap and pread/pwrite
  - saves calling system calls
- tuning parameters
  - bucket number, alignment
  - compression



### TreeDB: File Tree Database

### • B+ tree

- O(log N) time complexity
- custom comparison function

### page caching

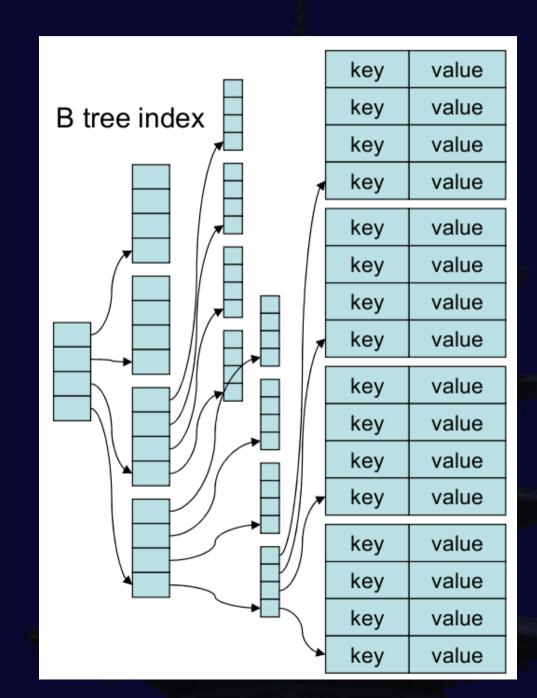
- separate LRU lists
- mid-point insertion

### stands on HashDB

- stores pages in HashDB
- succeeds time and space efficiency

### cursor

- jump, step\_back
- prefix/range matching



# HashDB vs. TreeDB

	HashDB (hash table)	TreeDB (B+ tree)		
time efficiency	faster, O(1)	slower, O(log N)		
space efficiency	larger, 16 bytes/record footprint	smaller, 3 bytes/record footprint		
concurrency	higher, locked by the record	lower, locked by the page		
cursor	undefined order	ordered by application logic		
	jump by full matching, unable to step back	jump by range matching, able to step back		
fragmentation	more incident	less incident		
random access	faster	slower		
sequential access	slower	faster		
transaction	faster	slower		
memory usage	smaller	larger		
compression	less efficient	more efficient		
tuning points	number of buckets, size of mapped memory	size of each page, capacity of the page cache		
typical use cases	session management of web service	job/message queue		
	user account database	sub index of relational database		
	document database	dictionary of words		
	access counter	inverted index of full-text search		
	cache of content management system	temporary storage of map-reduce		
	graph/text mining	archive of many small files		

# Directory Databases

- DirDB: directory hash database
  - respective files in a directory of the file system
  - suitable for large records
  - strongly depends on the file system performance
    - ReiserFS > EXT3 > EXT4 > EXT2 > XFS > NTFS
- ForestDB: directory tree database
  - B+ tree stands on DirDB
  - reduces the number of files and alleviates the load
  - the most scalable approach among all database types

# On-memory Databases

- ProtoDB: prototype database template
  - DB wrapper for STL map
  - any data structure compatible std::map are available
  - ProtoHashDB: alias of ProtoDB<std::unordered\_map>
  - ProtoTreeDB: alias of ProtoDB<std::map>
- StashDB: stash database
  - static hash table with packed serialization
  - much less memory usage than std::map and std::unordered\_map
- CacheDB: cache hash database
  - static hash table with doubly-linked list
  - constant memory usage
  - LRU (least recent used) records are removed
- GrassDB: cache tree database
  - B+ tree stands on CacheDB
  - much less memory usage than CacheDB

# performance characteristics

class name	persistence	algorithm	complexity	sequence	lock unit
ProtoHashDB	volatile	hash table	0(1)	undefined	whole (rwlock)
ProtoTreeDB	volatile	red black tree	O(log N)	lexical order	whole (rwlock)
StashDB	volatile	hash table	0(1)	undefined	record (rwlock)
CacheDB	volatile	hash table	0(1)	undefined	record (mutex)
GrassDB	volatile	B+ tree	O(log N)	custom order	page (rwlock)
HashDB	persistent	hash table	0(1)	undefined	record (rwlock)
TreeDB	persistent	B+ tree	O(log N)	custom order	page (rwlock)
DirDB	persistent	undefined	undefined	undefined	record (rwlock)
ForestDB	persistent	B+ tree	O(log N)	custom order	page (rwlock)

# Tuning Parameters

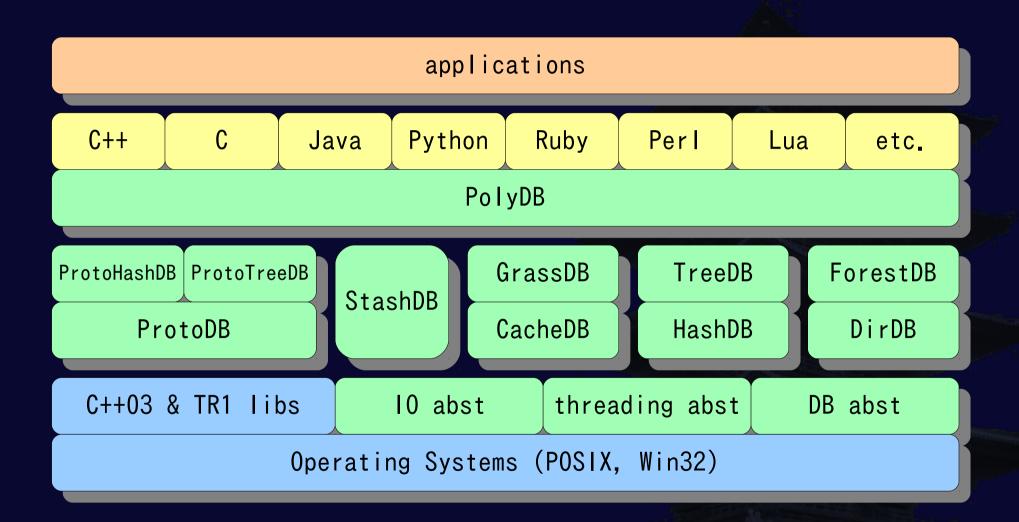
### Assumption

- stores 100 million records whose average size is 64 bytes
- pursues time efficiency first, and space efficiency second
- the machine has 16GB RAM
- HashDB: suited for usual cases
  - opts=TLINEAR, bnum=200M, msiz=12G
- TreeDB: suited for ordered access
  - opts=TLINEAR, bnum=10M, msiz=10G, pccap=2G
- StashDB: suited for fast on-memory associative array
  - bnum=100M
- CacheDB: suited for cache with LRU deletion
  - bnum=100M, capcnt=100M
- GrassDB: suited for memory saving on-memory associative array
  - bnum=10M, psiz=2048, pccap=2G
- DirDB and ForestDB: not suited for small records

# Class Hierarchy

- DB = interface of record operations
  - BasicDB = interface of storage operations, mix-in of utilities
    - ProtoHashDB, ProtoTreeDB, CacheDB, HashDB, TreeDB, ...
    - PolyDB
- PolyDB: polymorphic database
  - dynamic binding to concrete DB types
    - "factory method" and "strategy" patterns
  - the concrete type is determined when opening
    - naming convention
      - ProtoHashDB: "-", ProtoTreeDB: "+", CacheDB: "\*", GrassDB: "%"
      - HashDB: ".kch", TreeDB: ".kct", DirDB: ".kcd", ForestDB: ".kcf"

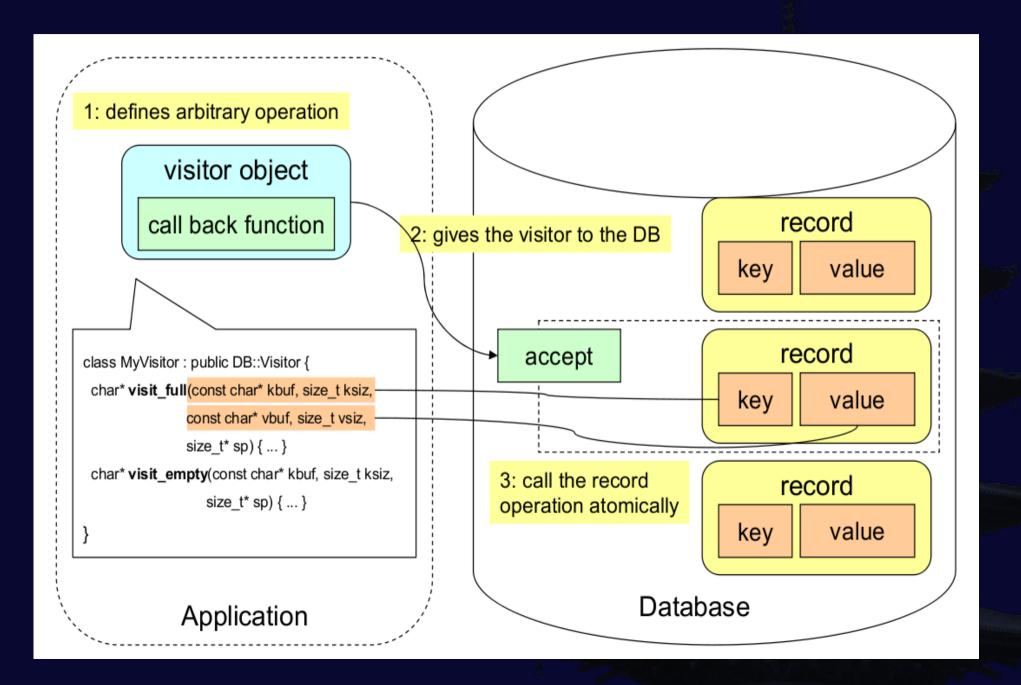
# Components



## Abstraction of KVS

- what is "key-value storage"?
  - each record consists of one key and one value
  - atomicity is assured for only one record
  - records are stored in persistent storage
- so, what?
  - every operation can be abstracted by the "visitor" pattern
  - the database accepts one visitor for a record at the same time
    - lets it read/write the record arbitrary
    - saves the computed value
- flexible and useful interface
  - provides the "DB::accept" method realizing whatever operations
  - "DB::set", "DB::get", "DB::remove", "DB::increment", etc. are built-in as wrappers of the "DB::accept"

# Visitor Interface



# Comparison with Tokyo Cabinet

### • pros

- space efficiency: smaller size of DB file
  - footprint/record: TC=22B -> KC=16B
- parallelism: higher performance in multi-thread environment
  - uses atomic operations such as CAS
- portability: non-POSIX platform support
  - supports Win32
- usability: object-oriented design
  - external cursor, generalization by the visitor pattern
  - new database types: StashDB, GrassDB, ForestDB
- robustness: auto transaction and auto recovery

### • cons

- time efficiency per thread: due to grained lock
- discarded features: fixed-length database, table database
- dependency on modern C++ implementation: jilted old environments

# Example Codes

```
#include <kcpolydb.h>
using namespace std;
using namespace kyotocabinet;
// main routine
int main(int argc, char** argv) {
  // create the database object
  PolyDB db;
  // open the database
  if (!db.open("casket.kch", PolyDB::OWRITER | PolyDB::OCREATE)) {
    cerr << "open error: " << db.error().name() << endl;</pre>
  // store records
 if (!db.set("foo", "hop") ||
!db.set("bar", "step") ||
      !db.set("baz", "jump")) {
    cerr << "set error: " << db.error().name() << endl;</pre>
  // retrieve a record
  string* value = db.get("foo");
  if (value) {
    cout << *value << endl;</pre>
    delete value;
  } else {
    cerr << "get error: " << db.error().name() << endl;</pre>
  // traverse records
  DB::Cursor* cur = db.cursor();
  cur->jump();
  pair<string, string>* rec;
  while ((rec = cur->get_pair(true)) != NULL) {
    cout << rec->first << ":" << rec->second << endl;</pre>
    delete rec;
  delete cur;
  // close the database
  if (!db.close()) {
    cerr << "close error: " << db.error().name() << endl;</pre>
  return 0;
```

```
#include <kcpolvdb.h>
using namespace std;
using namespace kyotocabinet;
// main routine
int main(int argc, char** argv) {
  // create the database object
  PolyDB db;
  // open the database
 if (!db.open("casket.kch", PolyDB::OREADER)) {
    cerr << "open error: " << db.error().name() << endl;</pre>
  // define the visitor
  class VisitorImpl : public DB::Visitor {
    // call back function for an existing record
    const char* visit full(const char* kbuf, size t ksiz,
                            const char* vbuf, size_t vsiz, size_t *sp) {
      cout << string(kbuf, ksiz) << ":" << string(vbuf, vsiz) << endl;</pre>
      return NOP;
    // call back function for an empty record space
    const char* visit_empty(const char* kbuf, size_t ksiz, size_t *sp) {
      cerr << string(kbuf, ksiz) << " is missing" << endl;</pre>
      return NOP;
  } visitor;
  // retrieve a record with visitor
 if (!db.accept("foo", 3, &visitor, false) ||
      !db.accept("dummy", 5, &visitor, false)) {
    cerr << "accept error: " << db.error().name() << endl;</pre>
  // traverse records with visitor
  if (!db.iterate(&visitor, false)) {
    cerr << "iterate error: " << db.error().name() << endl;</pre>
  // close the database
  if (!db.close()) {
   cerr << "close error: " << db.error().name() << endl;</pre>
  return 0;
```



- lightweight database server -

### **Features**

- persistent cache server based on Kyoto Cabinet
  - persistent associative array = key-value storage
  - client/server model = multi processes can access one database
- auto expiration
  - expiration time can be given to each record
    - records are removed implicitly after the expiration time
  - "GC cursor" eliminates expired regions gradually
- high concurrency and performance
  - resolves "c10k" problem with epoll/kqueue
  - 1M queries / 25 sec = 40,000 QPS or more
- high availability
  - hot backup, update logging, asynchronous replication
  - background snapshot for persistence of on-memory databases

### protocol based on HTTP

- available from most popular languages
- provides access library to encapsulate the protocol
- supports bulk operations also in an efficient binary protocol
- supports "pluggable server" and includes a memcached module

### various database schema

- using the polymorphic database API of Kyoto Cabinet
- almost all operations of Kyoto Cabinet are available
- supports "pluggable database"

### scripting extension

- embeds Lua processors to define arbitrary operations
- supports visitor, cursor, and transaction

### library of network utilities

- abstraction of socket, event notifier, and thread pool
- frameworks of TCP server, HTTP server, and RPC server

# Networking

- abstraction of system-dependent mechanisms
  - encapsulates threading mechanisms
    - thread, mutex, rwlock, spinlock, condition variable, TLS, and atomic integer
    - task queue with thread pool
  - encapsulates Berkeley socket and Win32 socket
    - client socket, server socket, and name resolver
  - encapsulates event notifiers
    - "epoll" on Linux, "kqueue" on BSD, "select" on other POSIX and Win32
    - substitutes for "libevent" or "libev"

### • TSV-RPC

- lightweight substitute for XML-RPC
- input data and output data are associative arrays of strings
  - serialized as tab separated values
- RESTful interface
  - GET for db::get, PUT for db::set, DELETE for db::remove

# Example Messages

### db::set over TSV-RPC

POST /rpc/set HTTP/1.1 Content-Length: 22

Content-Type: text/tab-separated-values

key japan value tokyo

HTTP/1.1 200 OK Content-Length: 0

### db::get over TSV-RPC

POST /rpc/get HTTP/1.1

Content-Length: 10

Content-Type: text/tab-separated-values

key japan

HTTP/1.1 200 OK Content-Length: 12

Content-Type: text/tab-separated-values

value tokyo

### db::set over RESTful

PUT /japan HTTP/1.1 Content-Length: 5

Content-Type: application/octet-stream

tokyo

HTTP/1.1 201 Created Content-Length: 0

### db::get over RESTful

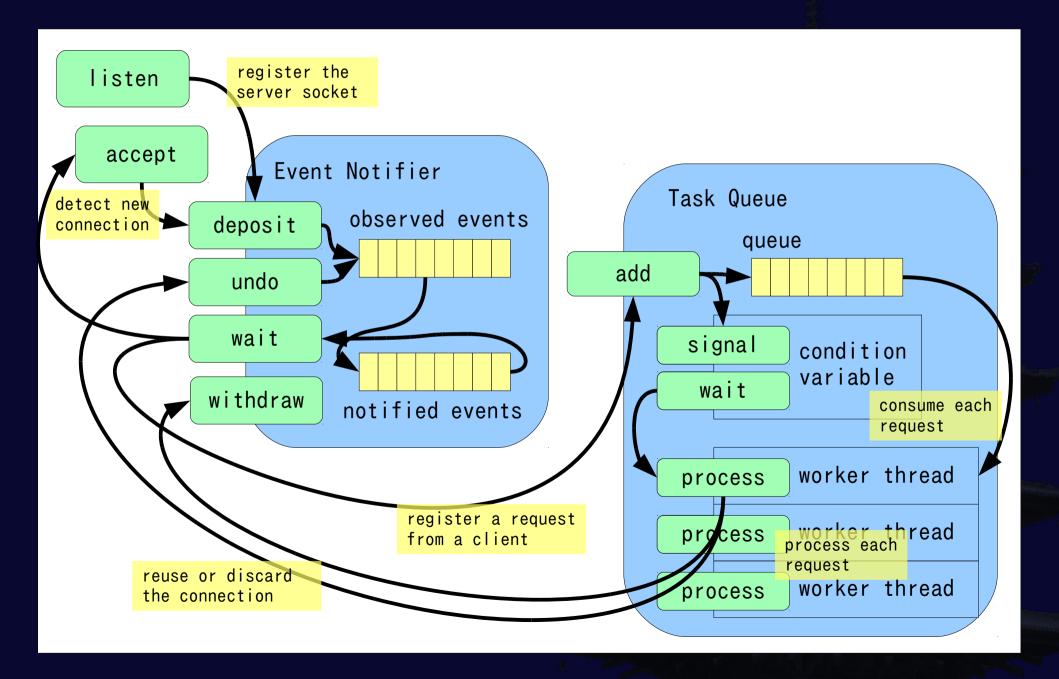
GET /japan HTTP/1.1

HTTP/1.1 200 OK Content-Length: 5

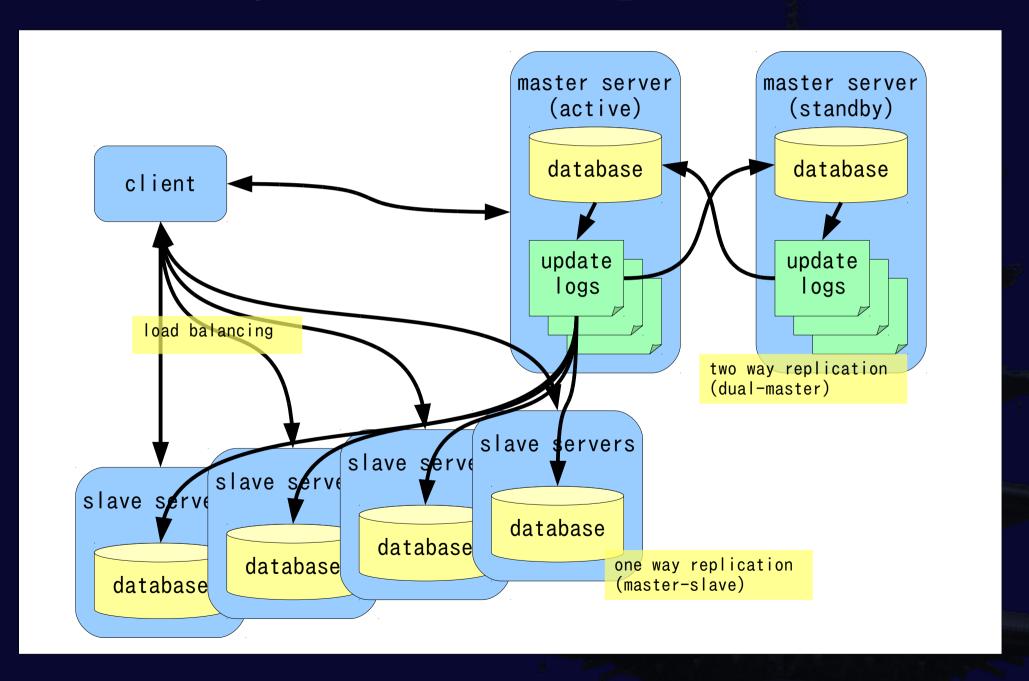
Content-Type: application/octet-stream

tokyo

# Threaded TCP Server

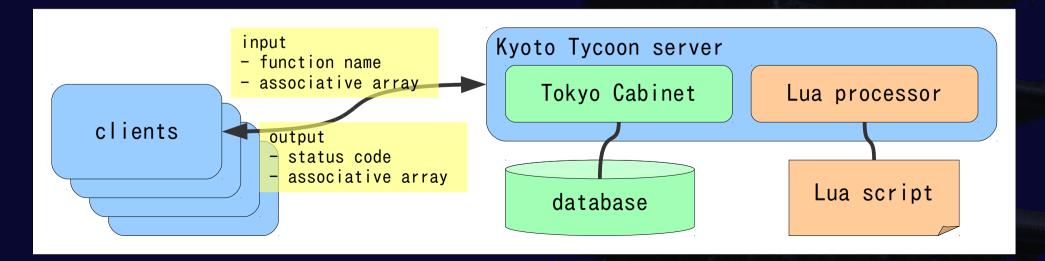


# Asynchronous Replication



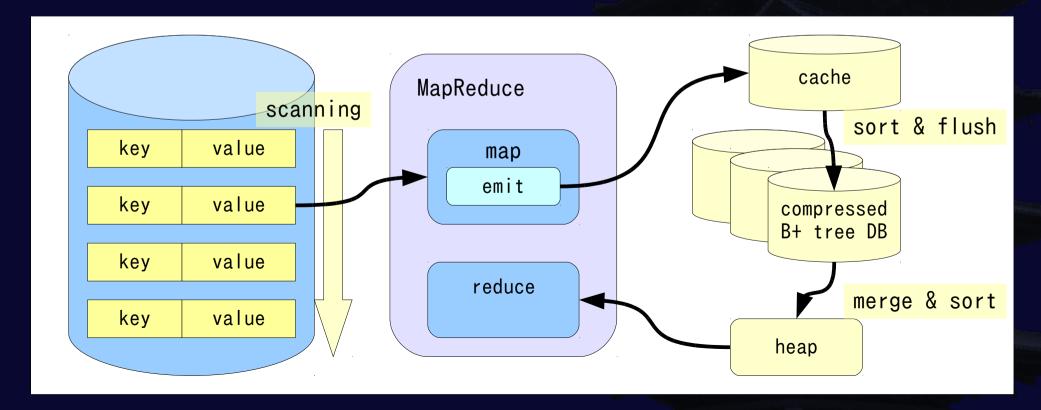
# Scripting Extension with Lua

- prominent time efficiency and space efficiency
  - suitable for embedding use
- extreme concurrency
  - each native thread has respective Lua instance
- poor functionality, compensated for by C functions
  - full set of the Lua binding of Kyoto Cabinet
  - pack, unpack, split, codec, and bit operation utilities



# Local MapReduce in Lua

- for on-demand aggregate calculation
  - the "map" and the "reduce" functions defined by users in Lua
- local = not distributed
  - fast cache and sort implementation in C++, no overhead for networking
  - space saving merge sort by compressed B+ trees



# Components

applications

database client

database server

Lua processor

TSV-RPC client

TSV-RPC server

HTTP client

HTTP server

client socket

threaded TCP server

networking utilities

server socket

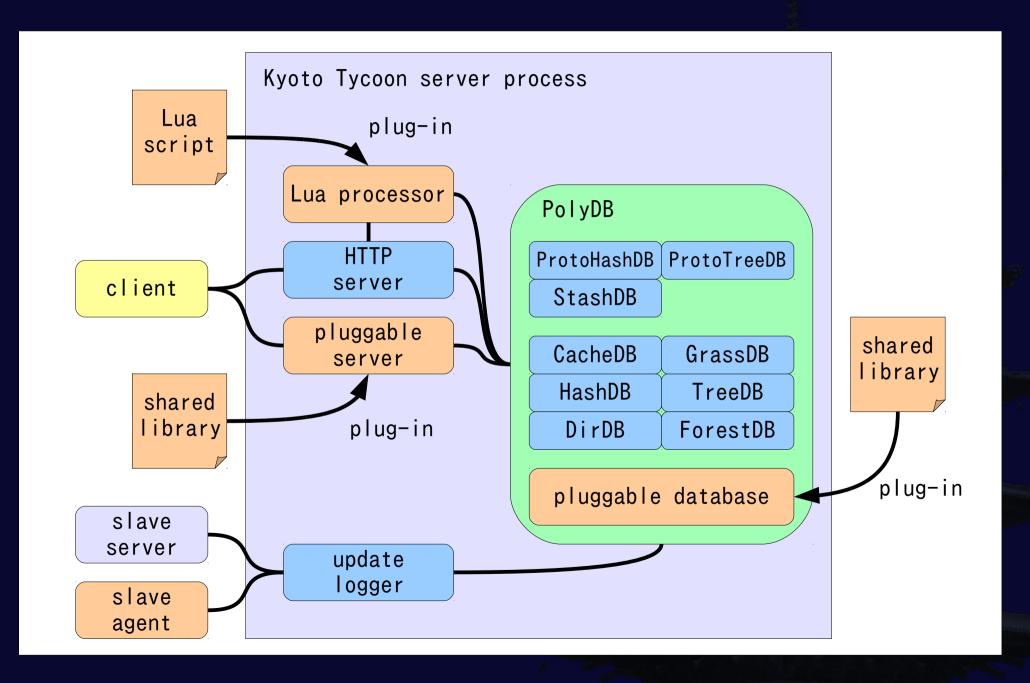
event notifier

Kyoto Cabinet (database, threading utilities)

Operating Systems (POSIX, Win32)

C++03 & TR1 libs

# pluggable modules



# Comparison with Tokyo Tyrant

### pros

- using Kyoto Cabinet
  - inherits all pros of Kyoto Cabinet against Tokyo Cabinet
- auto expiration
  - easy to use as a cache server
- multiple databases
  - one server can manage multiple database instances
- thorough layering
  - database server on TSV-RPC server on HTTP server on TCP server
  - easy to implement client libraries
- portability: non-POSIX platform support
  - supports Win32 (work-in-progress)

### cons

• discarded features: fixed-length database, table database

# Example Codes

```
#include <ktremotedb.h>
using namespace std;
using namespace kyototycoon;
// main routine
int main(int argc, char** argv) {
  // create the database object
  RemoteDB db;
  // open the database
  if (!db.open()) {
    cerr << "open error: " << db.error().name() << endl;</pre>
  // store records
 if (!db.set("foo", "hop") ||
    !db.set("bar", "step") ||
      !db.set("baz", "jump")) {
    cerr << "set error: " << db.error().name() << endl;</pre>
  // retrieve a record
  string* value = db.get("foo");
  if (value) {
    cout << *value << endl;</pre>
    delete value;
  } else {
    cerr << "get error: " << db.error().name() << endl;</pre>
  // traverse records
  RemoteDB::Cursor* cur = db.cursor();
  cur->jump();
  pair<string, string>* rec;
  while ((rec = cur->get_pair(NULL, true)) != NULL) {
    cout << rec->first << ":" << rec->second << endl;</pre>
    delete rec;
  delete cur;
  // close the database
 if (!db.close()) {
    cerr << "close error: " << db.error().name() << endl;</pre>
  return 0;
```

```
kt = kvototvcoon
db = kt.db
-- retrieve the value of a record
function get(inmap, outmap)
  local key = inmap.key
  if not key then
     return kt.RVEINVALID
  local value, xt = db:qet(key)
  if value then
     outmap.value = value
     outmap.xt = xt
  else
     local err = db:error()
     if err:code() == kt.Error.NOREC then
         return kt.RVELOGIC
     end
     return kt.RVEINTERNAL
  return kt.RVSUCCESS
end
-- list all records
function list(inmap, outmap)
  local cur = db:cursor()
  cur:jump()
  while true do
     local key, value, xt = cur:get(true)
     if not key then break end
     outmap.kev = value
  return kt.RVSUCCESS
end
```

# Other Kyoto Products?

now, planning...

- Kyoto Diaspora?
  - inverted index using Kyoto Cabinet
- Kyoto Prominence?
  - content management system using Kyoto Cabinet

maintainability is our paramount concern...

FAL Labs http://fallabs.com/mailto:info@fallabs.com/

# 長長人 キャビネット 8 EiB